

Decision Log

1. How the Policy Was Structured as Input to the AI

To make the Falls Management Policy usable within the system, I converted the policy requirements into a structured rule-based format inside `src/data/policy.js`.

The policy is divided into three sections corresponding to the three documentation stages:

- Incident Report (Day 1)
- Continuation Note (Day 2)
- Closure Note (Day 3)

Each requirement contains two parts:

1. **Extraction Schema**

- Defines the information that needs to be identified from the nurse's note.
- The AI is responsible only for extracting these pieces of information as boolean values.

2. **Evaluation Logic**

- A JavaScript function evaluates the extracted information and determines whether the requirement is:
 - Complete
 - Missing
 - Vague

I intentionally separated information extraction from compliance evaluation. This approach keeps the compliance logic transparent and easy to maintain because any future policy changes can be handled by updating the rule definitions rather than rewriting prompts.

2. One Suggestion from the AI That I Modified

While testing the extraction prompts, I noticed that the AI tended to classify many unclear statements as completely missing.

For example, if a nurse wrote:

“Seems okay this morning.”

the AI often treated this as if mobility had not been mentioned at all.

After reviewing several examples, I felt that this did not reflect how nurses commonly document observations in practice. To improve accuracy, I introduced two levels of interpretation:

- A broader check to determine whether the topic was mentioned.
- A stricter check to determine whether the required clinical detail was explicitly documented.

Using this approach, a statement such as “Seems okay this morning” can be recognized as an observation while still being flagged as lacking sufficient clinical detail. This helped produce more meaningful compliance results.

3. On The Most Important Decision That Required Human Judgment

One of the most challenging design decisions was deciding how to distinguish between **Vague** and **Missing** documentation.

Consider the note:

“Seems okay this morning.”

There are two possible interpretations:

- The nurse has not provided enough information to satisfy the requirement.
- The nurse has documented an observation, but the documentation lacks the detail needed for compliance.

This is not something the AI can reliably decide on its own because it depends on how the organisation wants to interpret documentation quality.

I chose to treat these situations as **Vague** rather than **Missing**. My reasoning was that the nurse has clearly made an observation, but the documentation does not provide enough evidence to confirm compliance with the policy.

To support this decision, I implemented separate checks for:

- Whether a topic was mentioned.
- Whether the required clinical information was explicitly confirmed.

This allows the system to provide more balanced feedback and reduces the risk of incorrectly flagging genuine observations as completely missing.